

Environment Setup to Access Hosted LLMs from Groq AI Provider:

- Groq is an AI provider that hosts and serves Large Language Models through REST APIs.
- Instead of downloading and running LLMs locally, developers can use Groq's hosted models by creating an account, generating an API key, and making API requests.
- Groq does not build all the models it serves.
- Just like Ollama hosts Open-Weight LLMs on your local machine, Groq hosts multiple LLMs in the cloud and provides high-performance access to them through REST APIs.
- Groq also offers a free usage tier (subject to current limits), making it a great choice for learning, testing, and developing AI applications.
- **Environment Setup to Access Hosted LLMs from Groq**
 - Create Groq account using : Console URL: <https://console.groq.com>
 - Generate API Key: <https://console.groq.com/keys>
 - Available Models:
 - <https://console.groq.com/docs/models>
 - <https://api.groq.com/openai/v1/models>
 - LLM Call (REST API): POST <https://api.groq.com/openai/v1/chat/completions>
 - Note: we have to pass body with required model name and request/prompt
 - View Rate Limits: <https://console.groq.com/settings/limits>
- Note: Do not store API keys directly in your application code. Instead, store them as OS environment variables or securely manage them using Kubernetes Secrets.